



Devexperts



dxFeed

From Demo to Production

A Year of Agentic AI in Financial Markets

Anton Antonov, Head of AI & Quant, dxFeed

Munich-London, September 2026

About dxFeed

dxFeed is a subsidiary of [Devexperts](#), with the primary focus of delivering financial markets information to buy-side and sell-side institutions and individuals, specifically to traders, data analysts, quants and portfolio managers.

6M

End-Users
through direct
and B2B2C
relationships

200,000

Simultaneous streaming
clients

3.5M

Instruments
(Equities, Futures,
Options, Indices,
FX, Crypto, etc.)

>100

B2B clients around
the world

What we shipped this year




Product	Headline
Grenadier-2	Capable of streaming full US equity universe @ 4,000 RPS; production latency p50 51 ms
GDELT News -> SummerFox	Licence-free news source, live in SummerFox; includes stocks, commodities, crypto
dxMCP	Market data + Grenadier + SummerFox on one MCP surface

The model stopped being the interesting part

Everything that makes AI work in production now lives in the harness around it.

Three proofs from one year of shipping:


1. a gateway that turns products into tools,
2. a pipeline that buys model judgement by the call,
3. and a product with no LLM in it at all.




dxMCP: what's on the surface


Model Context Protocol: the open standard for connecting AI applications to external systems.

Layer	On the surface today
Market data	Candles / OHLC, trades & quotes, instrument profiles (IPF), fundamentals, news
Screening	Options and underlying scanners - 400+ datapoints
Products as tools	Grenadier anomaly scores, SummerFox per-symbol and portfolio summaries
Consumers	AI agents and human operators, e.g. any MCP client on a desktop Products: DXcharts, DXtrade AI agent, SummerFox





Products as tools, not just apps



Decision rule #1: when MCP is the right answer

Use it when


- The answer already exists somewhere authoritative
- The model just needs to fetch it correctly
- More than one consumer will want the same data

Don't use

- As a substitute for a calculation engine
- When the consumer is one fixed pipeline – just call the API
- To deal with a data-quality problem upstream

SummerFox: the market report, written by agents


- Ask about a stock or a portfolio, get a structured report in seconds
- Every number & data point comes from dxFeed: candles, news, fundamentals, instrument profiles
- One report, any language
- Per-symbol & portfolio views
- Now with licence-free financial news on Google's GDELT GKG (BigQuery public dataset)



dxFeed SummerFox in Action

What is the high-level outlook?

What data-driven insights are there?



A headline is not a ticker

What the text says	The trap	What gets stored
"Marvell to join S&P 500"	The louder name is not the subject	stock:MRVL primary index:SPX secondary
"[...] Source: Nasdaq", "A quiet day on Nasdaq"	The org name produces bogus NDAQ matches	nothing
"COMP rises 2%" with no class prefix	Compass Inc. or the Nasdaq Composite? The string cannot tell you	nothing, not guessed


Extracting symbol is a judgement – exactly where an LLM is both necessary and dangerous.

Still no silver bullet



Text volume: ~330 tokens / article	LLM cost: \$0.4 – \$8 per Mtok	Runtime: 6 s/call	Precision (article -> ticker)
All articles: ~49 K/day	\$6.5 – \$130/day	~8 hrs/day	27%

Four layers, each with one job



Text volume: ~330 tokens / article	LLM cost: \$0.4 – \$8 per Mtok	Runtime: 6 s/call	Precision (article -> ticker)
All articles: ~49 K/day	\$6.5 – \$130/day	~8 hrs/day	27%
Gated: 1350/day	<\$4/day	14 min/day	68%


Decision rule #2: where the model belongs

1. Put the model where **judgement is genuinely required.**
2. Put **deterministic code** on both sides of it.
3. Make something **other than the model** decide what gets stored.

This results in predictable costs, runtime, and precision.

Sometimes the harness contains no model at all

*Grenadier: a 0-100 anomaly score for the order book of every US-listed stock, once a second, from one in-house model.
No LLM anywhere in it!*







The needle in the haystack problem: rare & fuzzily defined events

Rule-based or feature engineered conditions:

- must be manually constructed or selected
- must be calibrated and recalibrated, may get outdated quickly

The bet: one universal model, or a zoo



The key risk:

"Transfer learning fails: a universal model that covers a wide instrument universe is unfeasible."

Fallback plan: a per-symbol model zoo.


- One unsupervised autoencoder, 20 GB of proprietary L2 order-book data
- Generalises to new instruments without retraining

Target vs delivered

	Target	Delivered
Coverage	2,600+ US stocks	Full US equity universe
Throughput	2,600 RPS (1/symbol/sec)	4,000 RPS
Latency (same region)	Sub-second	p50 51 ms p99 223 ms
Running cost	Fixed budget per month	~1/5 of the budget ~1/20 for a representative subset

dxFeed Grenadier in Action

Large volume concentration far from the spread



Decision rule #3: when to train your own model

	Use an LLM when... (SummerFox)	Train your own when... (Grenadier)
The input	is language: headlines, filings, a question a human typed	is a numeric signal: 20 GB of order-book snapshots
"Correct" means	a judgement a human expert would agree with	a statistic: reconstruction error against what normal looks like
The knowledge	lives in the world: companies, events, context	lives only in your data: there's no LLM that ever saw it
Volume & latency	thousands of calls a day, seconds are fine	thousands of inferences a second, forever, with sub-second latency (no LLM does that at any price!)
Economics	pay per call because each call is worth it	pay per day/month because load is tested and predictable

Six evals, three buckets

Type	Our layer	What it catches
Code	Data: assertions vs source of truth	Wrong values
Code	Tool: trace replay against golden traces (SummerFox)	Right answer, wrong call
Code	Config: label menus test-pinned to the prompt	Config drift
Code	Model: precision/recall on a held-out labelled set (Grenadier)	Detection quality
Model-graded	Agent: faithfulness, no fabrication (Inspect, Deepeval)	Plausible fiction
Human	Sampled review of production traces, back into the golden set	Everything we didn't think of

Bonus: quality control and AI economics

Noise floor: ~40% of articles come back different on identical input

Each prompt edit, no matter how trivial, is a new deployment

Multi-vendor fallback helps deal with outages and model degradations

No reproducibility: measure old-vs-old before you measure old-vs-new

Approach your AI-related costs on a per-call, not per-month, basis

Token accounting must be built into the audit trail

The blueprint

Does the answer already exist somewhere authoritative?



MCP *fetch it correctly; never compute it*

Is judgement genuinely required?



A model – gated approach *deterministic code on both sides, something else decides what is stored*

Narrow task, high volume, your own data?



Train your own model *no LLM does 2,600 inferences a second at any price*

Key call: establish key metrics and estimate design decisions before you start scaling.

Key Takeaways

- The model is a component. The harness is the product.
- Provide facts as tools via MCP.
- Rely on LLM judgement in a gated model.
- Form numeric, high-volume, high-throughput signals in your own specialised model.
- Price and measure every call before you pay for it, or the pilot never becomes a product.
- Sometimes the right harness contains no LLM at all.

dxFeed SummerFox: AI-driven market reporting



<https://dxfeed.com/solutions/summerfox-market-intelligence-plus-ai/>

dxFeed Grenadier: L2-based anomaly detection



<https://dxfeed.com/solutions/dxfeed-grenadier/>

Press release: <https://dxfeed.com/dxfeed-grenadier-recognized-as-best-ai-based-solution/>

Thank you!

Anton Antonov



antonov@dxfeed.com



[linkedin.com/in/tonytonov](https://www.linkedin.com/in/tonytonov)